Just IT
B2W group
Apprenticeships | Training | Recruitment

Data Technician
Week 3
Database Design & SQL

Name: Al Amin

Course Date: From 10/02/2025 To 13/02/2025

Table of contents

Day 1: Task 1

Please research and complete the below questions relating to key concepts of databases.

<table>
<tr><td>What is a primary key?</td><td>A primary key is a field (or a combination of fields) in a database table that uniquely identifies each record in that table.

Key characteristics of a primary key:

1. Uniqueness: Each value in the primary key column must be unique. No two records can have the same primary key value.
2. Non-nullability: Every record must have a value for the primary key. It cannot be null.
3. Immutability: The value of a primary key should not change once it is set, as it’s used to identify the record.
4. Efficiency: Primary keys are often indexed to make querying and searching for records faster.

For example, in a Students table, a StudentID column might serve as the primary key, where each student gets a unique ID to distinguish them from others.</td></tr>
<tr><td>How does this differ from a secondary key?</td><td>The main difference between a primary key and a secondary key (often referred to as a foreign key or alternate key, depending on context) lies in their role and usage within a database.
Primary Key:
1. Uniqueness: A primary key uniquely identifies each record in a table.
2. Mandatory: It cannot have null values.
3. Single per Table: There can only be one primary key in a table, which may consist of one or multiple columns (composite key).
4. Identification: It’s used for identifying and accessing individual rows in a table efficiently.
5. Indexing: It is often automatically indexed to improve search performance.</td></tr>
</table>

	Secondary Key (Alternate Key / Foreign Key): 1. **Non-unique**: A secondary key does not have to be unique. It may contain duplicate values. 2. **Optional**: Unlike the primary key, secondary keys may allow null values, depending on the context (e.g., a foreign key referencing another table). 3. **Multiple per Table**: A table can have multiple secondary keys. These could be other columns that also allow for efficient searching or indexing, but they do not guarantee uniqueness. 4. **Purpose**: A secondary key is typically used for **searching** or **filtering** data, but it doesn't guarantee uniqueness. For example, a secondary key can be used to search records based on a column that might not necessarily be unique (like a last name column in a table of users). 5. **Foreign Key**: In relational databases, a foreign key is a type of secondary key that establishes a relationship between two tables. It is a field (or combination of fields) in one table that points to the primary key in another table. **Example:** Consider a **Customers** table: • **Primary Key**: customer ID (each customer has a unique customer ID) • **Secondary Key**: email (multiple customers might share the same email domain, and it might not be unique in the database) If we have an **Orders** table: • **Foreign Key (Secondary Key)**: customer ID in the Orders table references the customer ID in the **Customers** table, linking orders to the specific customer. In short: • **Primary Key** = Uniquely identifies records in a table. • **Secondary Key** = Used for querying or establishing relationships but does not guarantee uniqueness.
How are primary and foreign keys related?	Primary and foreign keys are both crucial concepts in relational databases that help maintain data integrity and establish relationships between tables.

	1. **Primary Key**:
	  ◦ A **primary key** is a unique identifier for a record in a database table. Each table can have only one primary key.
	  ◦ It ensures that each record in the table is unique and can be identified without ambiguity.
	  ◦ It cannot contain NULL values, and all values must be unique.
	2. **Foreign Key**:
	  ◦ A **foreign key** is a column (or set of columns) in one table that refers to the primary key in another table.
	  ◦ It creates a relationship between two tables, ensuring that the value in the foreign key column matches a valid primary key in the referenced table, or it can be NULL.
	  ◦ A foreign key enforces referential integrity, meaning it ensures that you cannot insert a value in the foreign key column that does not exist in the primary key column of the referenced table.
	Relationship Between Primary and Foreign Keys:
	• The primary key in **Table A** is referenced by the foreign key in **Table B**. This establishes a relationship where **Table B** depends on **Table A**.
	Example:
	  ◦ **Table A**: Employees (with EmployeeID as the primary key)
	  ◦ **Table B**: Orders (with EmployeeID as a foreign key pointing to EmployeeID in the Employees table)
	• This relationship allows us to relate records in **Table B** (Orders) to specific records in **Table A** (Employees).
	In summary, the **primary key** uniquely identifies records in one table, while the **foreign key** in another table references that primary key, creating a link between the two tables.

Provide a real-world example of a one-to-one relationship	A **one-to-one** relationship in a database occurs when a single record in one table is associated with only one record in another table. A real-world example of this could be: **Real-World Example: Passport and Citizen**

	Imagine a system where each **citizen** can only have one **passport** and each **passport** can only be assigned to one **citizen**. This represents a **one-to-one** relationship.
	• **Citizen Table**:
	  ○ CitizenID (Primary Key)
	  ○ Name
	  ○ DateOfBirth
	• **Passport Table**:
	  ○ PassportID (Primary Key)
	  ○ CitizenID (Foreign Key referencing Citizen Table)
	  ○ PassportNumber
	  ○ DateIssued
	In this case:
	• Each **Citizen** has exactly one **Passport**.
	• Each **Passport** is assigned to exactly one **Citizen**.
	Characteristics of the One-to-One Relationship:
	1. **Primary Key** in the **Citizen Table** (CitizenID) is related to a **foreign key** in the **Passport Table** (CitizenID), ensuring the unique association.
	2. If a new **Passport** is issued, it is assigned to exactly one **Citizen**, and vice versa.
	3. You could enforce referential integrity by ensuring that every passport record has a corresponding citizen record.
	This example is a simple case of a **one-to-one** relationship, where both entities (citizen and passport) are linked uniquely with each other.

Provide a real-world example of a one-to-many relationship	A **one-to-many** relationship occurs when a single record in one table is associated with multiple records in another table. Here's a real-world example:
	Real-World Example: Department and Employees
	Imagine a company where each **department** can have multiple **employees**, but each **employee** works in only one **department**. This represents a **one-to-many** relationship.
	• **Department Table**:
	  ○ DepartmentID (Primary Key)
	  ○ DepartmentName
	• **Employee Table**:
	  ○ EmployeeID (Primary Key)

	◦ EmployeeName ◦ DepartmentID (Foreign Key referencing Department Table) In this case: • A **Department** can have multiple **Employees**. • Each **Employee** belongs to one and only one **Department**. **Characteristics of the One-to-Many Relationship:** 1. **DepartmentID** in the **Department Table** is referenced by a **foreign key** in the **Employee Table** (DepartmentID), linking each employee to a specific department. 2. A **Department** can have many **Employees**, but each **Employee** can only be associated with one **Department** at a time. 3. You could query the database to find all employees in a specific department, or find out which department a specific employee works in. For example: • **Department Table**: ◦ DepartmentID = 1, DepartmentName = "Engineering" ◦ DepartmentID = 2, DepartmentName = "Marketing" • **Employee Table**: ◦ EmployeeID = 101, EmployeeName = "John Doe", DepartmentID = 1 ◦ EmployeeID = 102, EmployeeName = "Jane Smith", DepartmentID = 1 ◦ EmployeeID = 103, EmployeeName = "Sam Brown", DepartmentID = 2 In this example, the **Engineering** department has multiple employees, but each employee is tied to just one department. This illustrates a typical **one-to-many** relationship.
Provide a real-world example of a many-to-many relationship	A **many-to-many** relationship occurs when multiple records in one table are associated with multiple records in another table. A common real-world example of a **many-to-many** relationship is the **students and courses** scenario in a university system. **Real-World Example: Students and Courses** In a university system, each **student** can enrol in multiple **courses**, and each **course** can have multiple **students** enrolled. This represents a **many-to-many** relationship.

- **Student Table**:
    - StudentID (Primary Key)
    - StudentName
    - Major
- **Course Table**:
    - CourseID (Primary Key)
    - CourseName
    - Credits

However, to model the **many-to-many** relationship, you would introduce a **junction table** (also called a **many-to-many relationship table**), which links the students to the courses. This junction table holds references (foreign keys) to both the **Student** and **Course** tables.

- **StudentCourse Table** (Junction Table):
    - StudentID (Foreign Key referencing the Student Table)
    - CourseID (Foreign Key referencing the Course Table)

Example Data:

- **Student Table**:
    - StudentID = 1, StudentName = "Alice", Major = "Computer Science"
    - StudentID = 2, StudentName = "Bob", Major = "Mathematics"
- **Course Table**:
    - CourseID = 101, CourseName = "Data Structures", Credits = 3
    - CourseID = 102, CourseName = "Calculus", Credits = 4
- **StudentCourse Table** (Junction Table):
    - StudentID = 1, CourseID = 101 (Alice is enrolled in Data Structures)
    - StudentID = 2, CourseID = 101 (Bob is enrolled in Data Structures)
    - StudentID = 1, CourseID = 102 (Alice is enrolled in Calculus)

Characteristics of the Many-to-Many Relationship:

1. A **Student** can be enrolled in multiple **Courses** (e.g., Alice is enrolled in both "Data Structures" and "Calculus").
2. A **Course** can have multiple **Students** enrolled (e.g., both Alice and Bob are enrolled in "Data Structures").

	3. The **StudentCourse Table** acts as a bridge to facilitate the relationship between **Students** and **Courses**. **Benefits of a Many-to-Many Relationship:** • It allows flexible enrolment where students can enrol in multiple courses, and each course can have many students. • You can easily query the database to find all courses a student is taking or find all students enrolled in a specific course. This example clearly illustrates the **many-to-many** relationship between **students** and **courses**.

Day 1: Task 2

Please research and complete the below questions relating to key concepts of databases.

What is the difference between a relational and non-relational database?	The main difference between a **relational database** and a **non-relational database** lies in how data is stored, structured, and accessed. Here's a breakdown of the key differences:
	1. Data Structure
	• **Relational Database (RDBMS)**:
	  ○ **Data is stored in tables** (also known as relations) consisting of rows and columns. Each table has a defined schema (structure), with columns representing different attributes of the data.
	  ○ Data in relational databases follows a strict **schema** (predefined structure), meaning the data types and relationships between tables are fixed.
	  ○ Examples: **MySQL**, **PostgreSQL**, **Oracle**, **SQL Server**.
	• **Non-Relational Database (NoSQL)**:
	  ○ **Data is stored in various formats**, such as key-value pairs, documents, graphs, or wide-column stores.

- Non-relational databases are generally **schema-less** or have a flexible schema, allowing data to be stored in a more flexible or dynamic format.
- Examples: **MongoDB** (document-based), **Cassandra** (wide-column store), **Redis** (key-value store), **Neo4j** (graph database).

2. Schema Flexibility

- **Relational Database**:
  - Data must adhere to a **strict schema** defined at the time of database creation.
  - Each row in a table must follow the same structure, meaning any changes to the schema (e.g., adding new columns or modifying data types) typically require database migrations or updates.
- **Non-Relational Database**:
  - These databases are **schema-less** or have flexible schemas, allowing different records (documents or entries) to have different structures.
  - You can add or remove fields without impacting other records, making it easier to adapt to evolving data requirements.

3. Data Integrity and Relationships

- **Relational Database**:
  - **Relationships** between tables are maintained using **foreign keys**, enforcing referential integrity and ensuring consistency between related data.
  - **ACID (Atomicity, Consistency, Isolation, Durability)** properties are strongly enforced to guarantee reliable transactions and data integrity.
- **Non-Relational Database**:

- Many non-relational databases do not have built-in support for complex relationships between data, although some (like graph databases) do.
- These databases tend to follow **eventual consistency** rather than strict ACID compliance. They prioritize **availability** and **partition tolerance** (as per the CAP theorem) over strict consistency in some use cases.

4. Scalability

- **Relational Database**:
  - Relational databases are generally **vertically scalable**. This means that to scale them, you need to upgrade the hardware (e.g., more CPU, RAM) on a single server.
  - Scaling horizontally (across multiple servers) is possible but can be more complex due to the need for maintaining consistency between servers.
- **Non-Relational Database**:
  - Non-relational databases are often **designed to scale horizontally**, meaning they can distribute data across multiple servers (or clusters) more easily.
  - This makes them more suitable for applications requiring high availability, large amounts of unstructured data, and flexible scalability.

5. Query Language

- **Relational Database**:
  - Uses **SQL (Structured Query Language)** for querying and managing data. SQL is a standardized language for defining and manipulating relational data.
- **Non-Relational Database**:
  - Non-relational databases often have their own query languages or interfaces. For example:
    - **MongoDB** uses a JSON-like query language.

- **Cassandra** uses CQL (Cassandra Query Language), similar to SQL but adapted for wide-column stores.
- **Redis** uses simple commands for key-value operations.

6. Use Cases

- **Relational Database**:
  - Well-suited for applications that require structured data and strong relationships between entities (e.g., **financial systems**, **customer management systems**, **enterprise applications**).
  - Ideal when you need to enforce **data integrity** and use complex **joins** and **queries**.
- **Non-Relational Database**:
  - Better for use cases that require **flexibility**, **scalability**, and handling of large amounts of **unstructured** or **semi-structured** data (e.g., **big data applications**, **real-time analytics**, **social media platforms**, **content management systems**).
  - Great when the data model is **dynamic** or when the relationships between data entities are **not rigid**.

Summary Table

Feature	**Relational Database (RDBMS)**	**Non-Relational Database (NoSQL)**
Data Structure	Tables with rows and columns (fixed schema)	Various formats (documents, key-value, graphs, wide-column)
Schema	Fixed schema, strict structure	Flexible or schema-less
Relationships	Strong relationships with foreign keys	Limited relationships (depending on type, e.g., graph DBs)
Data Integrity	ACID properties (strong consistency)	Eventual consistency (can vary based on use case)

Scalability	Vertical scaling (requires more powerful hardware)	Horizontal scaling (across many servers/clusters)
Query Language	SQL (Structured Query Language)	Varies by type (e.g., MongoDB uses a query language, Cassandra uses CQL)
Use Cases	Structured data, complex queries, strong data integrity	Unstructured/semi-structured data, scalability, flexibility

Conclusion:

- **Relational databases** are ideal for structured data, complex relationships, and strict data integrity needs.
- **Non-relational databases** excel in scalability, flexibility, and handling unstructured or semi-structured data, with use cases like big data, real-time applications, and rapid development environments.

What type of data would benefit off the non-relational model?

Why?

Non-relational databases (NoSQL) are particularly well-suited for handling certain types of data and use cases where traditional relational databases might not be as efficient. Here are some types of data that would benefit from a non-relational model, along with reasons why:

1. Unstructured or Semi-Structured Data

- **Examples**: Text files, social media posts, emails, logs, images, audio, video.
- **Why**: Non-relational databases can store and manage data without a fixed schema. For instance, a **document-based NoSQL database** (like **MongoDB**) allows you to store data in formats like JSON, which can have nested structures and different fields for each record. This flexibility makes it easy to handle data that doesn't fit neatly into rows and columns, such as multimedia files or unstructured content.

2. Big Data and High-Volume Data

- **Examples**: Clickstream data, IoT sensor data, web analytics, customer interactions, social media activity.

- **Why**: NoSQL databases are typically **designed to scale horizontally**, meaning they can spread data across multiple servers or clusters. This ability to scale out is crucial for managing vast amounts of data that would overwhelm a single server in a relational database. For example, **Apache Cassandra** and **HBase** are designed to handle large-scale, high-velocity data and support fast read and write operations across distributed systems.

3. Real-Time Data

- **Examples**: Real-time analytics, financial transaction data, live streaming data, gaming data, stock market feeds.
- **Why**: Many non-relational databases (like **Redis**, which is an in-memory key-value store) are optimized for extremely fast data retrieval, allowing them to efficiently process high-velocity, real-time data. They can provide sub-millisecond latency and scale quickly to support real-time applications, unlike relational databases that can struggle with performance under similar conditions.

4. Flexible or Evolving Data Models

- **Examples**: Data with frequent schema changes, rapidly changing data structures, or data without a fixed model (such as agile software development environments).
- **Why**: Non-relational databases are **schema-less** or support flexible schemas, which means you can easily add new attributes or change data structures without needing to perform database migrations or altering the entire schema. This makes them ideal for use cases where the data model is constantly evolving, such as **content management systems** or **product catalogs** that need to handle new attributes as requirements change.

5. Graph Data and Relationships

- **Examples**: Social networks, recommendation engines, fraud detection systems, network topologies.
- **Why**: **Graph databases** (e.g., **Neo4j**) are a type of non-relational database optimized for handling relationships between entities. Graph databases store data as nodes (entities) and edges (relationships), which makes them highly effective for querying complex relationships, like finding the shortest path between two people in a social

network or discovering recommendations based on user behaviour.

6. Distributed Data with High Availability

- **Examples**: Global web applications, multi-region e-commerce platforms, content delivery networks (CDNs).
- **Why**: Non-relational databases are often built to be **distributed** and can easily be deployed across multiple regions or cloud environments, making them highly available and fault-tolerant. They support the **CAP theorem** (Consistency, Availability, Partition tolerance), ensuring that the system can continue operating even when parts of it go down or become unreachable. This makes NoSQL databases ideal for applications that require **high availability** and the ability to handle failures gracefully, such as **online retailers** or **social media platforms**.

7. Data with No Clear Relationship Between Entities

- **Examples**: Sensor data, log entries, metadata, or loosely related user data.
- **Why**: Non-relational databases, especially **key-value stores** (e.g., **Redis**, **DynamoDB**), can store data with no predefined relationships between entries. These databases work well for scenarios where you need to store data that doesn't have complex relationships or need to perform operations like fast lookups based on a key. This makes them efficient for things like caching or storing session data.

8. Time-Series Data

- **Examples**: Monitoring data, sensor data, event logging, weather data, financial market data.
- **Why**: Non-relational databases like **InfluxDB** and **TimescaleDB** are optimized for time-series data. They can handle massive amounts of sequential data, making them ideal for storing time-based information where each data point is associated with a timestamp. Non-relational databases allow for fast, high-throughput writing of data, and querying historical trends is efficient as well.

Why Non-Relational Databases Are Ideal for These Types of Data:

1. **Scalability**: Non-relational databases are often distributed and horizontally scalable, which means they can easily grow to accommodate massive datasets across many machines. This is essential when handling big data or high-velocity data.
2. **Flexibility**: Non-relational databases allow you to store data in a variety of formats without a predefined schema. This makes them ideal for dynamic, unstructured, or semi-structured data.
3. **Performance**: Many NoSQL databases are optimized for fast read and write operations, particularly when handling large volumes of data. For instance, key-value stores can retrieve data almost instantly when using a key, making them perfect for use cases like caching or session storage.
4. **Adaptability to Changing Requirements**: Non-relational databases can evolve alongside the application, allowing schema-less storage. This is especially useful for environments where requirements shift rapidly or where data attributes are added or changed frequently.

Conclusion:

Non-relational databases are most beneficial when dealing with **unstructured or semi-structured data**, **large volumes of data**, **real-time data**, **high scalability requirements**, or **dynamic and evolving data models**. They're an excellent choice for applications like **social networks**, **big data analytics**, **real-time processing**, and **content management systems**, where the flexibility, speed, and ability to scale are more critical than strict schema enforcement or relational data integrity

Day 3: Task 1

Please research the below 'JOIN' types, explain what they are and provide an example of the types of data it would be used on.

Self-join

A **self join** is a type of join where a table is joined with itself. It's useful when you need to compare rows within the same table. This is typically done by treating the same table as two different "tables" by using aliases for each instance of the table. A self join is often used to find relationships within the same set of data.

Example:

Consider a **"Employees"** table that contains information about employees and their managers. The table might have the following columns:

- EmployeeID
- EmployeeName
- ManagerID (this is a foreign key that references the EmployeeID of the manager)

Sample Data:

EmployeeID	EmployeeName	ManagerID
1	John	NULL
2	Alice	1
3	Bob	1
4	Charlie	2
5	Dave	2

In this example:

- John is the manager of Alice and Bob.
- Alice is the manager of Charlie and Dave.

Self Join Query:

To list all employees along with their managers' names, you can use a **self join**.

SELECT E.EmployeeName AS Employee, M.EmployeeName AS Manager

FROM Employees E

```
LEFT JOIN Employees M
ON E.ManagerID = M.EmployeeID;
```

Result:

Employee	Manager
John	NULL
Alice	John
Bob	John
Charlie	Alice
Dave	Alice

In the query:

- E is the alias for the employees.
- M is the alias for their managers (which is the same table as the employees).
- The LEFT JOIN is used to ensure that if an employee doesn't have a manager (like John), they still appear in the result with a NULL for their manager.

Types of Data for Self Joins:

- **Hierarchical data**: Such as employees and managers, or parent-child relationships in a family tree.
- **Comparing values within the same table**: Like finding pairs of products that have the same price.
- **Relational data**: Where there is a need to compare one entity's attributes with another of the same type, for example comparing people in the same database for mutual connections or recommendations.

Self joins are particularly powerful when you need to analyse hierarchical relationships or compare different rows of data from the same source.

Right join	A **RIGHT JOIN** (or **RIGHT OUTER JOIN**) in SQL returns all records from the right table and the matched records from the left table. If there is no match, the result is NULL on the left side. This join is particularly useful when you want to ensure that all records from the right table are included, regardless of whether they have matching records in the left table.

Syntax:

```
SELECT columns
FROM table1
RIGHT JOIN table2
ON table1.column = table2.column;
```

Example:

Consider two tables:

1. **Customers** - contains customer information.
2. **Orders** - contains order information, where each order is linked to a customer.

Customers Table:

CustomerID	CustomerName
1	John
2	Alice
3	Bob

Orders Table:

OrderID	CustomerID	OrderAmount
101	1	500
102	2	300
103	1	200
104	4	150

Notice that the Orders table includes an order from a customer with CustomerID = 4, but there is no matching customer in the **Customers** table for this order.

Right Join Query:

If we want to list all orders, including those from customers who do not exist in the **Customers** table (like customer 4), we can use a **RIGHT JOIN**:

```
SELECT Orders.OrderID, Customers.CustomerName,
Orders.OrderAmount
FROM Customers
```

```
RIGHT JOIN Orders
ON Customers.CustomerID = Orders.CustomerID;
```

Result:

OrderID	CustomerName	OrderAmount
101	John	500
102	Alice	300
103	John	200
104	NULL	150

- For the order with OrderID = 104, there is no corresponding CustomerID = 4 in the **Customers** table, so the CustomerName is NULL.
- All records from the **Orders** table are included, even though one of them doesn't have a match in the **Customers** table.

Types of Data for Right Joins:

- **Records with optional relationships**: For example, if you want to display all orders, even if some of the customers no longer exist or are missing in the **Customers** table.
- **Sales or transactions data**: Where you need to ensure that all transactions are displayed, even if the customer information is missing or incomplete.
- **Incomplete or sparse relationships**: If there are some entities in the "right" table that may not have corresponding data in the "left" table (like orders without customers or employees without department assignments).

When to Use a RIGHT JOIN:

- **Ensuring full coverage of the right table's data**: If you have data that might be missing from the left table, but you want to retain all records from the right table.
- **Dealing with optional relationships**: If there's a situation where the right table's data could exist without corresponding data in the left table, like in the case of unmatched orders or transactions.

In contrast to a **LEFT JOIN**, which ensures that all rows from the left table appear, the **RIGHT JOIN** ensures that all rows from the right table are returned, making it useful in scenarios where the focus is on preserving the right table's completeness.

Full join

A **FULL JOIN** (or **FULL OUTER JOIN**) in SQL returns **all records** when there is a match in **either** the left table or the right table. If there is no match, the result will contain NULL values for missing matches from either table.

In other words, it combines the results of both **LEFT JOIN** and **RIGHT JOIN**, returning all rows from both tables, with NULL where there is no match.

Syntax:

```
SELECT columns
FROM table1
FULL JOIN table2
ON table1.column = table2.column;
```

Example:

Let's consider two tables:

1. **Customers** - contains customer information.
2. **Orders** - contains order information, where each order is linked to a customer.

Customers Table:

CustomerID	CustomerName
1	John
2	Alice
3	Bob
5	Dave

Orders Table:

OrderID	CustomerID	OrderAmount
101	1	500
102	2	300
103	1	200
104	4	150

In this example:

- Customers 1, 2, and 3 have orders, but Customer 5 does not have any orders.
- There is an order for Customer 4, but no corresponding customer with CustomerID = 4 in the **Customers** table.

Full Join Query:

To list all customers and their corresponding orders (if any), and all orders with their corresponding customers (if any), we can use a **FULL JOIN**:

```
SELECT Customers.CustomerName, Orders.OrderID,
Orders.OrderAmount
FROM Customers
FULL JOIN Orders
ON Customers.CustomerID = Orders.CustomerID;
```

Result:

CustomerName	OrderID	OrderAmount
John	101	500
John	103	200
Alice	102	300
Bob	NULL	NULL
Dave	NULL	NULL
NULL	104	150

Explanation:

- **John** has two orders (101 and 103), so both are shown with his name.
- **Alice** has one order (102), so it is shown with her name.
- **Bob** and **Dave** are customers, but they have no orders, so they are shown with NULL values for OrderID and OrderAmount.
- The order with OrderID = 104 exists for a customer (CustomerID = 4), but no matching customer exists in the **Customers** table, so the customer name is NULL for that order.

Types of Data for Full Joins:

- **Combining two datasets where one or both can have missing or incomplete relationships**: For example, when merging records from two different tables where one table may have some entries without corresponding entries in the other table.
- **Handling missing data**: When you're working with incomplete or missing relationships, like a list of customers and orders, where some customers have no orders, and some orders have no customer information.

	• **Data reconciliation**: Full joins are helpful in scenarios where you need to merge two data sets that might have non-matching entries, ensuring that all records are included from both tables. **When to Use a Full Join:** • **When you want to include all data from both tables**: Even when there are no matching records. • **For reporting or analysis where you need to see all possible records**: For example, listing customers with and without orders, or orders with and without customers. • **Data merging or cleaning**: When you want to combine two datasets and ensure that no data is left out, even if some data doesn't match. **Full Join vs. Left Join and Right Join:** • **FULL JOIN**: Includes all records from both tables, with NULL for non-matching rows. • **LEFT JOIN**: Includes all records from the left table, with NULL for non-matching rows from the right table. • **RIGHT JOIN**: Includes all records from the right table, with NULL for non-matching rows from the left table. The **FULL JOIN** is especially useful when you want to preserve all data, regardless of whether or not there are matches between the tables.
Inner join	An **INNER JOIN** in SQL returns only the rows where there is a **match** between the two tables based on the condition specified in the ON clause. If no match is found, the row is **excluded** from the result. **Key Points:** • **Only matched records are included** in the result. • It eliminates rows from both tables that don't meet the join condition. **Syntax:** SELECT columns FROM table1 INNER JOIN table2 ON table1.column = table2.column; **Example:** Let's consider two tables: 1. **Customers** - contains customer information.

2. **Orders** - contains order information, where each order is linked to a customer.

Customers Table:

CustomerID	CustomerName
1	John
2	Alice
3	Bob
5	Dave

Orders Table:

OrderID	CustomerID	OrderAmount
101	1	500
102	2	300
103	1	200
104	4	150

In this example:

- **Customers** 1, 2, and 3 have orders.
- **Customer** 5 has no orders.
- There is an order with OrderID = 104, but there is no customer with CustomerID = 4.

Inner Join Query:

To find the customers who have placed orders, we can use an **INNER JOIN**:

```
SELECT Customers.CustomerName, Orders.OrderID,
Orders.OrderAmount
FROM Customers
INNER JOIN Orders
ON Customers.CustomerID = Orders.CustomerID;
```

Result:

CustomerName	OrderID	OrderAmount
John	101	500
John	103	200
Alice	102	300

Explanation:

- Only the customers who have placed orders are included in the result.

	Bob and **Dave** are excluded because Bob has no orders, and Dave is not in the **Orders** table. - The order with OrderID = 104 is excluded because there is no matching CustomerID = 4 in the **Customers** table. **Types of Data for Inner Joins:** **Related data**: Inner joins are typically used when you want to find records in one table that are related to records in another table. **Transactional data**: For example, connecting customers with the orders they've placed, or linking employees with the projects they're working on. **Filtering out unmatched records**: If you need to exclude rows without matches from either table. **When to Use an Inner Join:** **When you're only interested in rows that have matching data in both tables**: For example, you might want to list only customers who have placed orders, excluding those who haven't. **To filter out incomplete or missing relationships**: If you only want the data where a valid relationship exists between the two tables, like matching products and sales or users and login records. **Inner Join vs. Other Joins:** **INNER JOIN**: Only returns rows where there is a match in both tables. **LEFT JOIN**: Returns all rows from the left table, and matching rows from the right table, with NULL for non-matching rows from the right table. **RIGHT JOIN**: Returns all rows from the right table, and matching rows from the left table, with NULL for non-matching rows from the left table. **FULL JOIN**: Returns all rows from both tables, with NULL for non-matching rows from either table. The **INNER JOIN** is a good choice when you want to combine data from two related tables but only want to include the rows where both tables have corresponding entries.
Cross join	A **CROSS JOIN** in SQL returns the **Cartesian product** of two tables, meaning it returns every possible combination of rows between the two tables. If table A has m rows and table B has n rows, a

CROSS JOIN will return m * n rows. It does not require a condition to join the tables, unlike other types of joins.

Key Points:

- The **CROSS JOIN** does not require a condition (like ON in other joins).
- It produces a **Cartesian product**—each row from the first table is paired with every row from the second table.
- The result can be quite large, especially if the tables have many rows.

Syntax:

```
SELECT columns
FROM table1
CROSS JOIN table2;
```

Example:

Let's use two tables:

1. **Colors** - contains a list of colors.
2. **Sizes** - contains a list of sizes.

Colors Table:

Color
Red
Green
Blue

Sizes Table:

Size
Small
Medium
Large

Cross Join Query:

To create all possible combinations of **Colors** and **Sizes**, we can use a **CROSS JOIN**:

```
SELECT Colors.Color, Sizes.Size
FROM Colors
CROSS JOIN Sizes;
```

Result:

Color	Size
Red	Small
Red	Medium

Red Large

Green Small

Green Medium

Green Large

Blue Small

Blue Medium

Blue Large

Explanation:

- The result contains every possible combination of a **Color** and a **Size**.
- If there are 3 colors and 3 sizes, the **CROSS JOIN** will produce 9 combinations (3 colors * 3 sizes).

Types of Data for Cross Joins:

- **Creating combinations of different attributes**: For example, generating a list of all possible combinations of product attributes (e.g., color and size, as shown above) when those attributes are stored in separate tables.
- **Generating test data**: Cross joins can be used to create a large number of combinations of data for testing purposes. For instance, generating a list of all possible combinations of dates and products for a sales report.
- **Combinatorial data**: When you need to explore every possible combination of two sets of data.

When to Use a Cross Join:

- **When you need all possible combinations** of two datasets. For example, to list all possible product configurations based on different attributes (like color and size).
- **When you want to generate combinations for scenarios like product bundles**, promotions, or options.
- **To produce test data or simulate combinations of events** that could happen based on two independent lists.

Cross Join vs. Other Joins:

- **CROSS JOIN**: Returns the Cartesian product of the two tables, without any condition, and can result in a very large result set.
- **INNER JOIN**, **LEFT JOIN**, **RIGHT JOIN**, **FULL JOIN**: These types of joins require a condition (ON clause) and return only the

	rows where there is a match or inclusion based on the specified condition, rather than all combinations. **Example Use Case in a Business Context:** Imagine you're running a store and want to list all combinations of products, colors, and sizes available. You could use a **CROSS JOIN** to combine a table of products with tables of available colors and sizes, helping you to generate a list of all product variations.
Left join	A **LEFT JOIN** (also known as a **LEFT OUTER JOIN**) in SQL returns **all rows** from the **left table** and the **matching rows** from the **right table**. If there is no match in the right table, the result will include NULL values for the columns from the right table. **Key Points:** • All rows from the left table will be included in the result, whether or not there is a match in the right table. • If no match exists, the result will contain NULL values for the right table's columns. **Syntax:** SELECT columns FROM table1 LEFT JOIN table2 ON table1.column = table2.column; **Example:** Let's consider two tables: 1. **Customers** - contains customer information. 2. **Orders** - contains order information, where each order is linked to a customer. **Customers Table:** <table><tr><th>CustomerID</th><th>CustomerName</th></tr><tr><td>1</td><td>John</td></tr><tr><td>2</td><td>Alice</td></tr><tr><td>3</td><td>Bob</td></tr><tr><td>5</td><td>Dave</td></tr></table>**Orders Table:** <table><tr><th>OrderID</th><th>CustomerID</th><th>OrderAmount</th></tr><tr><td>101</td><td>1</td><td>500</td></tr><tr><td>102</td><td>2</td><td>300</td></tr><tr><td>103</td><td>1</td><td>200</td></tr></table>

104	6	150

In this example:

- **Customers 1, 2, and 3** have orders.
- **Customer 5** has no orders.
- There is an order with OrderID = 104 for a customer (CustomerID = 6), but this customer does not exist in the **Customers** table.

Left Join Query:

To list all customers and their corresponding orders, but ensure that even customers without orders are included, we can use a **LEFT JOIN**:

```
SELECT Customers.CustomerName, Orders.OrderID,
Orders.OrderAmount
FROM Customers
LEFT JOIN Orders
ON Customers.CustomerID = Orders.CustomerID;
```

Result:

CustomerName	OrderID	OrderAmount
John	101	500
John	103	200
Alice	102	300
Bob	NULL	NULL
Dave	NULL	NULL

Explanation:

- **John** and **Alice** have orders, so their order details are shown.
- **Bob** and **Dave** are customers but have no corresponding orders in the **Orders** table, so their OrderID and OrderAmount are shown as NULL.
- The order with OrderID = 104 is **not** included because there is no corresponding customer with CustomerID = 6 in the **Customers** table.

Types of Data for Left Joins:

- **Optional relationships**: If you have a table with mandatory data (e.g., customers) and a related table with optional data (e.g., orders, payments), you can use a **LEFT JOIN** to include all customers, even those who haven't placed orders.

- **Incomplete or missing relationships**: For example, when you want to list all employees (left table), but some employees have not been assigned to any projects (right table).
- **Default or fallback data**: If you want to ensure that all records from the left table are included in the result, even if there is no matching record in the right table.

When to Use a Left Join:

- **To ensure that all rows from the left table are included**: Even if there are no matching rows in the right table.
- **To identify missing or incomplete relationships**: For example, finding customers who haven't made any purchases or employees who haven't been assigned to projects.
- **For reports where the main focus is on the left table**: Such as a list of products, customers, or employees, where you want to include even those who have no associated data in another table.

Left Join vs. Other Joins:

- **LEFT JOIN**: Returns all rows from the left table, and matching rows from the right table (with NULL if no match exists).
- **INNER JOIN**: Only returns rows where there is a match between the two tables.
- **RIGHT JOIN**: Similar to a LEFT JOIN, but includes all rows from the right table and matching rows from the left.
- **FULL JOIN**: Returns all rows from both tables, with NULL where there is no match in either table.

In summary, the **LEFT JOIN** is useful when you need to make sure that all records from the left table are shown in the result, even if there is no corresponding record in the right table.

Day 4: Task 1: Written

In your groups, discuss and complete the below activity. You can either nominate one writer or split the elements between you. Everyone however must have the completed work below:

Imagine you have been hired by a small retail business that wants to streamline its operations by creating a new database system. This database will be used to manage inventory, sales, and customer information. The business is a small corner shop that sells a range of groceries and domestic products. It might help to picture your local convenience store and think of what they sell. They also have a loyalty program, which you will need to consider when deciding what tables to create.

Write a 500-word essay explaining the steps you would take to set up and create this database. Your essay should cover the following points:

1. ***Understanding the Business Requirements**:*
    a. *What kind of data will the database need to store?*
    b. *Who will be the users of the database, and what will they need to accomplish?*
2. ***Designing the Database Schema**:*
    a. *How would you structure the database tables to efficiently store inventory, sales, and customer information?*
    b. *What relationships between tables are necessary (e.g., how sales relate to inventory and customers)?*
3. ***Implementing the Database**:*
    a. *What SQL commands would you use to create the database and its tables?*
    b. *Provide examples of SQL statements for creating tables and defining relationships between them.*
4. ***Populating the Database**:*
    a. *How would you input initial data into the database? Give examples of SQL INSERT statements.*
5. ***Maintaining the Database**:*
    a. *What measures would you take to ensure the database remains accurate and up to date?*
    b. *How would you handle backups and data security?*

Your essay should include specific examples of SQL commands and explain why each step is necessary for creating a functional and efficient database for the retail business.

Please write your 500-word essay here

Designing a Database System for a Small Retail Business

Creating a database system for a small retail business requires careful planning to ensure that it meets the operational needs of managing inventory, sales, and customer information. The goal is to design a system that not only stores relevant data but also allows users to efficiently retrieve and update that data. Below, I outline the steps I would take to set up and create the database for this retail business.

1. Understanding the Business Requirements

To begin, it's essential to understand the specific data the database will need to store. In a retail context, the database should track inventory, sales transactions, customer information, and loyalty program participation. Specifically:

- **Inventory Data**: Information on products, including product name, price, quantity in stock, and product categories (e.g., groceries, household items).
- **Sales Data**: Sales transactions, including the date of sale, products sold, quantity sold, total price, and any applicable discounts.
- **Customer Data**: Customer names, contact details, and loyalty program status (e.g., points earned, membership level).
- **Loyalty Program Data**: The number of points accumulated by customers, rewards, and program status.

Users of the database would include store managers, cashiers, and inventory controllers. The store manager might need to generate reports on sales trends or inventory status, while cashiers will need to quickly process sales and update inventory. The database should facilitate these operations with minimal complexity.

2. Designing the Database Schema

To structure the database efficiently, I would create the following tables:

- **Products**: This table will store information about the items for sale, such as name, price, and quantity.

- Columns: ProductID (Primary Key), Name, Price, QuantityInStock, Category.

- **Sales**: This table tracks each sale made, linking products and customers.
  - Columns: SaleID (Primary Key), SaleDate, CustomerID (Foreign Key), TotalAmount.
- **SalesDetails**: This table captures individual items in each sale, linking sales to products.
  - Columns: SaleDetailID (Primary Key), SaleID (Foreign Key), ProductID (Foreign Key), QuantitySold, PriceEach.
- **Customers**: Contains information about each customer, including loyalty program data.
  - Columns: CustomerID (Primary Key), Name, Email, LoyaltyPoints.
- **LoyaltyPrograms**: This table stores information on customer loyalty programs, including points thresholds for rewards.
  - Columns: LoyaltyID (Primary Key), CustomerID (Foreign Key), PointsAccumulated.

The relationships between the tables are as follows:

- The **Sales** table is linked to **Customers** through the CustomerID, indicating which customer made the purchase.
- The **Sales** table is also related to **SalesDetails**, which lists each product sold in a transaction.
- **SalesDetails** connects to **Products**, identifying which products were sold.
- **Customers** are linked to **LoyaltyPrograms** to track their loyalty points.

3. Implementing the Database

To implement this database, I would use SQL to create the tables and define relationships. Below are examples of SQL statements to create these tables:

```
CREATE TABLE Products (
    ProductID INT PRIMARY KEY,
    Name VARCHAR(255),
    Price DECIMAL(10,2),
    QuantityInStock INT,
```

```sql
    Category VARCHAR(100)
);

CREATE TABLE Customers (
    CustomerID INT PRIMARY KEY,
    Name VARCHAR(255),
    Email VARCHAR(255),
    LoyaltyPoints INT
);

CREATE TABLE Sales (
    SaleID INT PRIMARY KEY,
    SaleDate DATE,
    CustomerID INT,
    TotalAmount DECIMAL(10,2),
    FOREIGN KEY (CustomerID) REFERENCES
Customers(CustomerID)
);

CREATE TABLE SalesDetails (
    SaleDetailID INT PRIMARY KEY,
    SaleID INT,
    ProductID INT,
    QuantitySold INT,
    PriceEach DECIMAL(10,2),
    FOREIGN KEY (SaleID) REFERENCES Sales(SaleID),
    FOREIGN KEY (ProductID) REFERENCES Products(ProductID)
);

CREATE TABLE LoyaltyPrograms (
    LoyaltyID INT PRIMARY KEY,
    CustomerID INT,
    PointsAccumulated INT,
    FOREIGN KEY (CustomerID) REFERENCES
Customers(CustomerID)
);
```

4. Populating the Database

To input initial data into the database, I would use SQL INSERT statements. Here are examples of how to populate each table with sample data:

```sql
INSERT INTO Products (ProductID, Name, Price, QuantityInStock, Category)
VALUES (1, 'Milk', 1.99, 50, 'Groceries'),
       (2, 'Toothpaste', 3.49, 30, 'Household');

INSERT INTO Customers (CustomerID, Name, Email, LoyaltyPoints)
VALUES (1, 'John Doe', 'john.doe@email.com', 100);

INSERT INTO Sales (SaleID, SaleDate, CustomerID, TotalAmount)
VALUES (1, '2025-02-12', 1, 5.98);

INSERT INTO SalesDetails (SaleDetailID, SaleID, ProductID, QuantitySold, PriceEach)
VALUES (1, 1, 1, 2, 1.99);

INSERT INTO LoyaltyPrograms (LoyaltyID, CustomerID, PointsAccumulated)
VALUES (1, 1, 100);
```

5. Maintaining the Database

To ensure the database remains accurate and up to date, several practices should be followed:

- **Regular Updates**: Inventory levels should be updated in real-time during sales transactions. This can be automated using triggers or within the application logic.
- **Backups**: Periodic database backups should be scheduled to protect against data loss. Automated backups can be configured, and data can be stored securely in a cloud environment or offsite storage.
- **Security Measures**: Ensuring data security is essential. This includes restricting access to sensitive customer information using proper authentication and authorisation, and encrypting sensitive data fields such as email addresses or loyalty points.

In conclusion, creating a database for a small retail business involves understanding the business's operational needs, designing an efficient schema, implementing it using SQL, and ensuring ongoing maintenance to keep the system up to date and secure. By following these steps, the business can streamline its operations and improve customer service through better data management.

Day 4: Task 2: SQL Practical

In your groups, work together to answer the below questions. It may be of benefit if one of you shares your screen with the group and as a team answer / take screen shots from there.

Setting up the database:

1. **Download world_db(1) [here]()**
2. **Follow each step to create your database [here]()**

For each question I would like to see both the syntax used and the output.

1. **Count Cities in USA:** *Scenario:* You've been tasked with conducting a demographic analysis of cities in the United States. Your first step is to determine the total number of cities within the country to provide a baseline for further analysis.

Syntax: SELECT COUNT(DISTINCT Name) AS 'Number of cities in USA' FROM city WHERE countrycode = 'USA';

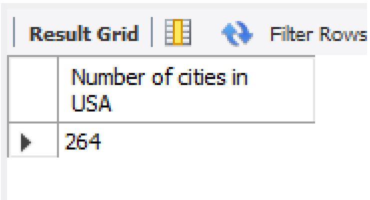

2. **Country with Highest Life Expectancy:** *Scenario:* As part of a global health initiative, you've been assigned to identify the country with the highest life expectancy. This information will be crucial for prioritising healthcare resources and interventions.

Syntax: Select name as 'Country Name', LifeExpectancy
from country
where LifeExpectancy=(select max(lifeexpectancy) from country);

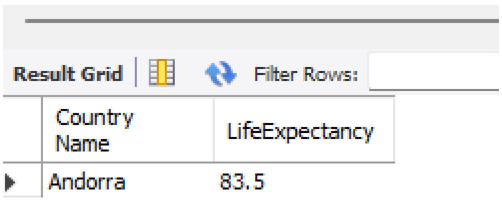

3. **"New Year Promotion: Featuring Cities with 'New :** *Scenario:* In anticipation of the upcoming New Year, your travel agency is gearing up for a special promotion featuring cities with names including the word 'New'. You're tasked with swiftly compiling a list of all cities from around the world. This curated selection will be essential in creating promotional materials and enticing travellers with exciting destinations to kick off the New Year in style.

Syntax: *Select Name from city where name like 'new %' ;*

Result Grid

Name
New Bombay
New Delhi
New York
New Orleans
New Haven
New Bedford

4. **Display Columns with Limit (First 10 Rows):** *Scenario:* You're tasked with providing a brief overview of the most populous cities in the world. To keep the report concise, you're instructed to list only the first 10 cities by population from the database.

Syntax:
```
Select name as 'City Name', Population from city
order by population Desc  limit 10;
```

Result Grid | Filter Rows:

City Name	Population
Mumbai (Bombay)	10500000
Seoul	9981619
São Paulo	9968485
Shanghai	9696300
Jakarta	9604900
Karachi	9269265
Istanbul	8787958
Ciudad de México	8591309
Moscow	8389200
New York	8008278

5. **Cities with Population Larger than 2,000,000:** *Scenario:* A real estate developer is interested in cities with substantial population sizes for potential investment opportunities. You're tasked with identifying cities from the database with populations exceeding 2 million to focus their research efforts.

Syntax: Select name as 'City Name', Population
from city where population>2000000 order by population desc ;

Result Grid | Filter Rows:

City Name	Population
Mumbai (Bombay)	10500000
Seoul	9981619
São Paulo	9968485
Shanghai	9696300
Jakarta	9604900
Karachi	9269265
Istanbul	8787958
Ciudad de México	8591309
Moscow	8389200
New York	8008278
Tokyo	7980230
Peking	7472000
London	7285000
Delhi	7206704
Cairo	6789479

city 35 ×

6. **Cities Beginning with 'Be' Prefix:** *Scenario:* A travel blogger is planning a series of articles featuring cities with unique names. You're tasked with compiling a list of cities from the database that start with the prefix 'Be' to assist in the blogger's content creation process.

Syntax: Select Name from city
where name like 'Be%' order by name Asc ;

Result Grid | Filter Rows:

Name
Beau Bassin-Rose Hill
Beaumont
Beawar
Béchar
Beerseba
Bei´an
Beihai
Beipiao
Beira
Beirut
Béjaïa
Bekasi
Belém
Belfast
Belford Roxo

city 38 ×

7. **Cities with Population Between 500,000-1,000,000:** *Scenario:* An urban planning committee needs to identify mid-sized cities suitable for infrastructure development projects. You're tasked with identifying cities with populations ranging between 500,000 and 1 million to inform their decision-making process.

Syntax: Select name as 'City Name', Population from city where population between 500000 and 1000000 order by population desc ;

Result Grid | Filter Rows:

City Name	Population
Amman	1000000
Mogadishu	997000
Volgograd	993400
Sendai	989975
Peshawar	988005
Baotou	980000
Adelaide	978100
Madurai	977856
Mekka	965700
Köln	962507
Managua	959000
Detroit	951270
Shenzhen	950500
Haora (H...	950435

city 39

8. **Display Cities Sorted by Name in Ascending Order:** *Scenario:* A geography teacher is preparing a lesson on alphabetical order using city names. You're tasked with providing a sorted list of cities from the database in ascending order by name to support the lesson plan.

Syntax: `Select * from city order by name asc;`

Result Grid | Filter Rows: | Edit: | Export/Import:

ID	Name	CountryCode	District	Population
698	[San Cristóbal de] la Laguna	ESP	Canary Islands	127945
20	´s-Hertogenbosch	NLD	Noord-Brabant	129170
670	A Coruña (La Coruña)	ESP	Galicia	243402
3097	Aachen	DEU	Nordrhein-Westfalen	243825
3318	Aalborg	DNK	Nordjylland	161161
2760	Aba	NGA	Imo & Abia	298900
1404	Abadan	IRN	Khuzestan	206073
395	Abaetetuba	BRA	Pará	111258
3683	Abakan	RUS	Hakassia	169200
1849	Abbotsford	CAN	British Colombia	105403
2747	Abeokuta	NGA	Ogun	427400
478	Aberdeen	GBR	Scotland	213070
3191	Abha	SAU	Asir	112300
2812	Abidjan	CIV	Abidjan	2500000
1703	Abiko	JPN	Chiba	126670

city 40

9. **Most Populated City:** *Scenario:* A real estate investment firm is interested in cities with significant population densities for potential development projects. You're tasked with identifying the most populated city from the database to guide their investment decisions and strategic planning.

Syntax: Select name as 'City', Population from city where population=(select max(population) from city);

Result Grid | Filter Rows:

	City	Population
▶	Mumbai (Bombay)	10500000

10. **City Name Frequency Analysis: Supporting Geography Education** *Scenario*: In a geography class, students are learning about the distribution of city names around the world. The teacher, in preparation for a lesson on city name frequencies, wants to provide students with a list of unique city names sorted alphabetically, along with their respective counts of occurrences in the database. You're tasked with this sorted list to support the geography teacher.

Syntax: SELECT name AS 'City', COUNT(*) AS 'Number_of_Times_City_Appears' FROM city GROUP BY name ORDER BY name ASC;

Result Grid | Filter Rows: | Export: | Wra

	City	Number_of_Times_City_Appears
▶	[San Cristóbal de] la Laguna	1
	´s-Hertogenbosch	1
	A Coruña (La Coruña)	1
	Aachen	1
	Aalborg	1
	Aba	1
	Abadan	1
	Abaetetuba	1
	Abakan	1
	Abbotsford	1
	Abeokuta	1
	Aberdeen	1
	Abha	1
	Abidjan	1
	Abiko	1

Result 45 ×

Syntax: SELECT name AS 'City', COUNT(*) AS 'Occurences_of_Name' FROM city GROUP BY name ORDER BY count(*) DESC;

	City	Occurences_of_Name
▶	San José	4
	Córdoba	3
	San Miguel	3
	San Fernando	3
	Hamilton	3
	La Paz	3
	Toledo	3
	Cambridge	3
	Springfield	3
	Richmond	3
	Valencia	3
	León	3
	Victoria	3
	Jining	2
	Kansas City	2

11. **City with the Lowest Population:** *Scenario:* A census bureau is conducting an analysis of urban population distribution. You're tasked with identifying the city with the lowest population from the database to provide a comprehensive overview of demographic trends.

Syntax: Select name as 'City', Population from city where population=(select min(population) from city);

Result Grid | Filter Rows

City	Population
Adamstown	42

12. **Country with Largest Population:** *Scenario:* A global economic research institute requires data on countries with the largest populations for a comprehensive analysis. You're tasked with identifying the country with the highest population from the database to provide valuable insights into demographic trends.

Syntax: Select name as 'Coutry', Population from country where population=(select max(population) from country);

Result Grid | Filter Rows:

Coutry	Population
China	1277558000

13. **Capital of Spain:** *Scenario:* A travel agency is organising tours across Europe and needs accurate information on capital cities. You're tasked with identifying the capital of Spain from the database to ensure itinerary accuracy and provide travellers with essential destination information.

Syntax: SELECT country.name AS 'Country', city.name AS 'Capital' FROM city INNER JOIN country ON city.ID = country.capital WHERE country.name = 'Spain';

Result Grid | Filter Rows:

Country	Capital
Spain	Madrid

14. **Cities in Europe:** *Scenario:* A European cultural exchange program is seeking to connect students with cities across the continent. You're tasked with compiling a list of cities located in Europe from the database to facilitate program planning and student engagement.

Syntax: SELECT country.name AS 'Country', city.name AS 'Capital'
FROM city INNER JOIN country ON city.ID = country.capital
WHERE country.Continent = 'Europe';

Result Grid | Filter Rows: | Export:

Country	Capital
Albania	Tirana
Andorra	Andorra la Vella
Austria	Wien
Belgium	Bruxelles [Brussel]
Bulgaria	Sofija
Bosnia and Herzegovina	Sarajevo
Belarus	Minsk
Switzerland	Bern
Czech Republic	Praha
Germany	Berlin
Denmark	København
Spain	Madrid
Estonia	Tallinn
Finland	Helsinki [Helsingf...
France	Paris

15. **Average Population by Country:** *Scenario:* A demographic research team is conducting a comparative analysis of population distributions across countries. You're tasked with calculating the average population for each country from the database to provide valuable insights into global population trends.

Syntax: Select avg(population) as 'Average Population' from country;

Average Population
25434098.1172

*Syntax: SELECT name AS Country, CONCAT(round(population * 100.0 / (SELECT SUM(population) FROM country), 2), '%') AS Percentage_of_Global_Population*
FROM country
*ORDER BY (population * 100.0 / (SELECT SUM(population) FROM country)) DESC*;

Country	Percentage_of_Global_Population
China	21.02%
India	16.68%
United States	4.58%
Indonesia	3.49%
Brazil	2.80%
Pakistan	2.57%
Russian Federation	2.42%
Bangladesh	2.12%
Japan	2.08%
Nigeria	1.83%
Mexico	1.63%
Germany	1.35%
Vietnam	1.31%
Philippines	1.25%
Egypt	1.13%

Result 83

16. **Capital Cities Population Comparison:** *Scenario:* A statistical analysis firm is examining population distributions between capital cities worldwide. You're tasked with comparing the populations of capital cities from different countries to identify trends and patterns in urban demographics.

Syntax: *SELECT Country.Name AS Country, City.Name AS City, City.Population*
FROM City INNER JOIN Country ON City.ID = Country.Capital
ORDER BY population DESC;

Result Grid | Filter Rows: | Export: | Wrap

Country	City	Population
South Korea	Seoul	9981619
Indonesia	Jakarta	9604900
Mexico	Ciudad de México	8591309
Russian Federation	Moscow	8389200
Japan	Tokyo	7980230
China	Peking	7472000
United Kingdom	London	7285000
Egypt	Cairo	6789479
Iran	Teheran	6758845
Peru	Lima	6464693
Thailand	Bangkok	6320174
Colombia	Santafé de Bogotá	6260862
Congo, The Demo...	Kinshasa	5064000
Chile	Santiago de Chile	4703954
Iraq	Baghdad	4336000

Result 60 ×

17. **Countries with Low Population Density:** *Scenario:* An agricultural research institute is studying countries with low population densities for potential agricultural development projects. You're tasked with identifying countries with sparse populations from the database to support the institute's research efforts.

Syntax: *SELECT name AS 'Country', population AS 'Population'*
FROM country where population!= 0 ORDER BY population ASC;

Result Grid | Filter Rows: | Export: | W

Country	Population
Pitcairn	50
Cocos (Keeling) Islands	600
Holy See (Vatican City State)	1000
Falkland Islands	2000
Norfolk Island	2000
Niue	2000
Tokelau	2000
Christmas Island	2500
Svalbard and Jan Mayen	3200
Saint Helena	6000
Saint Pierre and Miquelon	7000
Anguilla	8000
Montserrat	11000
Nauru	12000
Tuvalu	12000

country 66 ×

18. **Display Columns with Limit (Rows 31-40):** *Scenario:* A market research firm requires detailed information on cities beyond the top rankings for a comprehensive analysis. You're tasked with providing data on cities ranked between 31st and 40th by population to ensure a thorough understanding of urban demographics.

Syntax: SELECT ROW_NUMBER() OVER (ORDER BY City.Population DESC) AS 'Row_Number', City.Name AS City, Country.Name AS Country, City.Population FROM City INNER JOIN Country ON City.ID = Country.Capital ORDER BY City.Population DESC LIMIT 10 OFFSET 30;

Result Grid | Filter Rows: | Export:

Row_Number	City	Country	Population
31	Alger	Algeria	2168000
32	Paris	France	2125246
33	Toskent	Uzbekistan	2117500
34	Luanda	Angola	2022000
35	Bucuresti	Romania	2016131
36	Caracas	Venezuela	1975294
37	Brasília	Brazil	1969868
38	Budapest	Hungary	1811552
39	Baku	Azerbaijan	1787800
40	Kabul	Afghanistan	1780000

Course Notes

It is recommended to take notes from the course, use the space below to do so, or use the revision guide shared with the class:

We have included a range of additional links to further resources and information that you may find useful, these can be found within your revision guide.

END OF WORKBOOK

Please check through your work thoroughly before submitting and update the table of contents if required.

Please send your completed work booklet to your trainer.